

Chapter 5 - matrices and Linear Algebra

Solution Workbook for Student Practical Classes

(c) Dr Rob Collins 2023

2024-10-12

Table of contents

5	matrices and Linear Algebra	5 - 3
5.1	Introduction	5 - 3
5.2	Instructions for Students	5 - 3
5.3	Two Libraries for Linear Algebra - Sympy and Numpy	5 - 5
5.4	Algebraic Equations represented in matrices	5 - 5
5.5	matrices using the Sympy Library	5 - 5
5.5.1	Defining matrices in Sympy	5 - 5
5.5.2	Matrix Addition and Subtraction	5 - 6
5.5.3	Matrix multiplication by a constant	5 - 6
5.5.4	Multiplying pairs of matrices	5 - 7
5.5.5	Dot-Product (Inner Product) of Vectors	5 - 8
5.5.6	Cross-product of Vectors	5 - 8
5.5.7	Vector Norm - Length of a Vector	5 - 9
5.5.8	Determinant and Inverse of a Square Matrix	5 - 9
5.5.9	Other Matrix Operations from Linear Algebra	5 - 11
5.6	matrices in Numpy	5 - 13
5.6.1	Matrix Addition	5 - 14
5.6.2	Matrix Subtraction	5 - 15
5.6.3	Matrix multiplication by a Constant	5 - 15
5.6.4	Dot-product (Sometimes called ‘inner product’)	5 - 16
5.6.5	Matrix Division by a constant	5 - 17
5.6.6	Determinant of a matrix	5 - 18
5.6.7	Inverse of a matrix	5 - 18
5.6.8	Solutions to Linear Equations	5 - 19
5.6.9	Matrix (or vector) Norm	5 - 20

List of Figures

1	The 1853 attempt to solve ‘the matrix problem’. The so called ‘Steam Rivulet Engine’ was capable of solving large-scale, matrix calculations but unfortunately suffered from frequent gasket blow-outs and ultimately proved to be too dangerous for common use	5 - 4
---	---	-------

5 matrices and Linear Algebra

5.1 Introduction

This Machine Learning course is intended to be generally non-mathematical. Nevertheless, no advanced study of Machine Learning can be realistically considered complete without a good grounding in Probability and Linear Algebra.

For students who are less familiar with these subjects I recommend the following as useful references:

Strang, G. (2005) “Linear Algebra and its Applications”, TBS

Walpole, R.E., Myers, R.H., Myers, S.L. and Ye, K. (2012) “Probability and Statistics for Engineers and Scientists”, Prentice Hall

Deisenroth, M.P., Faisal, A.A. and Ong, C.S (2021) “Mathematics for Machine Learning”, Cambridge University Press

For many years Prof. Gilbert Strang presented a ‘famous’ course in Linear Algebra at MIT. The course is available free as part of MITs ‘Open Courseware’ programme. If you wish to become proficient in Machine Learning, that course is a strong recommendation: <https://ocw.mit.edu/courses/18-06-linear-algebra-spring-2010/> . My own experience was that both Strang’s book and the lecture course videos are easier to follow if you have access to a computer-tool to experiment with some of the Algebra. The SymPy tool used in this section provides exactly such a tool and may be a significant aid to study.

The following tutorial provides what I regard as an minimum for required for a first course in Machine Learning. The material in this tutorial will be a requirement in order to read and understand most text-books, papers and even websites in the field of Machine Learning.

5.2 Instructions for Students

This workbook is different to some of the others in this series. Unlike the other workbooks, I have provided all of the code you need. This is because the best way of describing the operations in each case is to simply show the code.

However, I would still encourage you to type in all of the code snippets and experiment with them. Modifying and experimenting with this code is an excellent way of becoming familiar with operations on matrices and Linear Algebra. Having completed this workbook you may find that the recommended course text book on Linear Algebra (Strang) is rather easier to follow.




Figure 1: The 1853 attempt to solve ‘the matrix problem’. The so called ‘Steam Rivulet Engine’ was capable of solving large-scale, matrix calculations but unfortunately suffered from frequent gasket blow-outs and ultimately proved to be too dangerous for common use

5.3 Two Libraries for Linear Algebra - Sympy and Numpy

This tutorial introduces two libraries for linear algebra + Sympy + Numpy

The first is a *symbolic* mathematics library - and this make it easier to read and understand linear algebra as it might appear in a mathematics text-book (e.g. Strang).

The second is a very fast and efficient library used for numerical computation. It implements many more linear algebra operations than Sympy and is much more widely used in Machine Learning.

So Sympy is the best tool to help you learn linear algebra whereas you are much more likely to use Numpy for real Machine Learning projects.

5.4 Algebraic Equations represented in matrices

Linear Algebra is concerned with the use of equations such as:

$$1x + 2y = 3$$

$$4x + 5y = 6$$

In the context of Machine Learning you will often see such equations represented in the form of matrices:

$$\begin{bmatrix} 1 & 2 \\ 4 & 5 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 3 \\ 6 \end{bmatrix}$$

At this point you should check your understanding of how the above equations are represented in matrix form

5.5 matrices using the Sympy Library

5.5.1 Defining matrices in Sympy

For example, we can represent the following two matrices:

$$x = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

.. and

$$y = \begin{bmatrix} 7 & 8 \\ 9 & 3 \\ 2 & 1 \end{bmatrix}$$

Let's start by using Sympy to define some matrices which we can use in later examples

```

1 from sympy import *
2
3 x = Matrix([[1,2,3], [4,5,6]])
4 x

```

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

```

1 y = Matrix([[7,8],[9,3],[2,1]])
2 y

```

$$\begin{bmatrix} 7 & 8 \\ 9 & 3 \\ 2 & 1 \end{bmatrix}$$

```

1 p = Matrix([[4,7,1],[4,3,2]])
2 p

```

$$\begin{bmatrix} 4 & 7 & 1 \\ 4 & 3 & 2 \end{bmatrix}$$

5.5.2 Matrix Addition and Subtraction

Two matrices can be added only if they are the same shape. Addition of matrices, simply means adding each element in one array to the corresponding element in the second array:

```

1 x + p

```

$$\begin{bmatrix} 5 & 9 & 4 \\ 8 & 8 & 8 \end{bmatrix}$$

The same goes for subtraction. matrices have to be the same shape, and they are subtracted on an element-by-element basis:

```

1 x - p

```

$$\begin{bmatrix} -3 & -5 & 2 \\ 0 & 2 & 4 \end{bmatrix}$$

5.5.3 Matrix multiplication by a constant

matrices can be multiplied by a constant. Each element is multiplied by the constant:

```
1 2 * x
```

$$\begin{bmatrix} 2 & 4 & 6 \\ 8 & 10 & 12 \end{bmatrix}$$

5.5.4 Multiplying pairs of matrices

matrices can be multiplied together but they have to be the correct shape. The number of columns in x has to match the number of rows in y. Notice also the order of multiplication and the shape of the result

Remember that x has a shape of 2x3:

```
1 x
```

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

y has a shape of 3x2

```
1 y
```

$$\begin{bmatrix} 7 & 8 \\ 9 & 3 \\ 2 & 1 \end{bmatrix}$$

So the result will be a 2x2 matrix:

```
1 x * y
```

$$\begin{bmatrix} 31 & 17 \\ 85 & 53 \end{bmatrix}$$

And the multiplication won't work with if the matrices are of the wrong shape:

```
1 k = Matrix([[1,2],[3,4]])
2 k
```

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

The next cell will give an error:

```
1 x * k
```

ShapeError: Matrix size mismatch: (2, 3) * (2, 2).

5.5.5 Dot-Product (Inner Product) of Vectors

It is frequently useful in Machine Learning to compute the 'dot-product' between two vectors. The dot-product measures the projection (shadow) of one vector onto another. The result is a scalar quantity.

```
1 v1 = Matrix([1,2,3,4,5])
2 v2 = Matrix([3,2,1,7,6])
3 v1
```

$$\begin{bmatrix} 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{bmatrix}$$

```
1 v2
```

$$\begin{bmatrix} 3 \\ 2 \\ 1 \\ 7 \\ 6 \end{bmatrix}$$

```
1 v1.dot(v2)
```

68

Which, because of symmetry is exactly the same as `v2.v1`:

```
1 v2.dot(v1)
```

68

5.5.6 Cross-product of Vectors

The other, common, vector operation is the cross-product. This essentially provides a measure of the area or volume 'enclosed' or 'marked out' by the vectors. It returns a vector result such that the magnitude of the vector is the area / volume and the direction of the vector is normal to the volume / area defined.

Sympy will only perform this operation with vectors of maximum length 3

```
1 v3 = Matrix([1,2,3])
2 v4 = Matrix([3,2,1])
3 v3
```


$$\begin{bmatrix} 1 \\ 2 \\ 3 \end{bmatrix}$$

```
1 v4
```

$$\begin{bmatrix} 3 \\ 2 \\ 1 \end{bmatrix}$$

```
1 v3.cross(v4)
```

$$\begin{bmatrix} -4 \\ 8 \\ -4 \end{bmatrix}$$

5.5.7 Vector Norm - Length of a Vector

In machine learning we will often want to compute the length of a vector. There are actually many (an infinite number!) of ways of doing this .. but the most common is computed in exactly the same way as in normal geometry: the square-root of the sum of the squares of the coordinates.

Interestingly, Sympy shows the result symbolically rather than as a real number .. but then again, that is exactly what Sympy was designed to do!

```
1 v1.norm()
```

$$\sqrt{55}$$

```
1 v2.norm()
```

$$3\sqrt{11}$$

5.5.8 Determinant and Inverse of a Square Matrix

If you are not familiar with the definition, use and computation of determinants then you should check out one of the references (such as Strang) for an explanation. Essentially, square matrices can represent a transformation in m-dimensional space. The determinant indicates the scaling of an object (represented by a matrix / vector) in that space from before to after the transformation. If the determinant is zero, that means that the object is squished to zero volume (length) by the transformation. One implication of this is that there is no inverse defined for a matrix with determinant of zero : once an object is squished flat we can't un-squish it again.

Sympy enables easy computation of the determinant of a square matrix.

```
1 A = Matrix([[1, 2, 3], [3, 6, 2], [2, 0, 1]] )
2 A
```

$$\begin{bmatrix} 1 & 2 & 3 \\ 3 & 6 & 2 \\ 2 & 0 & 1 \end{bmatrix}$$

```
1 A.det()
```

−28

A has a non-zero determinant and therefore the matrix has an inverse:

```
1 A.inv()
```

$$\begin{bmatrix} -\frac{3}{14} & \frac{1}{14} & \frac{1}{2} \\ -\frac{1}{28} & \frac{5}{28} & -\frac{1}{4} \\ \frac{3}{7} & -\frac{1}{7} & 0 \end{bmatrix}$$

However, if we started with a matrix with a determinant of zero:

```
1 Z = Matrix([[1,2,3],[4,5,6],[7,8,9]])
2 Z
```

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}$$

```
1 Z.det()
```

0

.. the inverse of Z is not defined (the following will show an error)

```
1 Z.inv()
```

NonInvertibleMatrixError: Matrix det == 0; not invertible.

5.5.9 Other Matrix Operations from Linear Algebra

We will often want to think about matrices as concise representations of a system of linear equations. In that case there are several operations that will help us think about and compute properties of that system:

- Determinant - (as above) - is the matrix invertible / system of equations ‘soluble’)
- Rank - How many rows / columns in the matrix are independent
- Column space - The sub-space which is ‘accessible’ (‘marked out’) by the vectors in the matrix (which may not be a full m-dimensional space, even though the matrix has dimension $m \times n$)
- Null space - For a matrix ‘A’ - the set of all vectors ‘x’ that satisfy $Ax = 0$

The last of these are somewhat advanced concepts. If you are not familiar with these, then you are directed towards Gilbert Strangs “Linear Algebra” for an excellent explanation.

Consider this matrix, representing three equations with 4 unknowns:

```
1 A = Matrix([[1,3,3,2], [2,6,9,7], [-1,-3,3,4]])
2 A
```

$$\begin{bmatrix} 1 & 3 & 3 & 2 \\ 2 & 6 & 9 & 7 \\ -1 & -3 & 3 & 4 \end{bmatrix}$$

But there are only actually 2 independent rows / columns here:

```
1 A.rank()
```

2

That is illustrated by using Gaussian Elimination to transform the matrix into reduced row echelon form:

```
1 U = A.rref()
2 U[0]
```

$$\begin{bmatrix} 1 & 3 & 0 & -1 \\ 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

(The slightly annoying ‘[0]’ here is because the sympy ‘rref’ function actually returns a vector result. The zeroth element of that vector is the reduced echelon form. The second element indicates the ‘pivot’ columns. The pivot columns are those with that mark out the column space of the matrix. Another way of saying this, is that they are the columns that have non-zero numbers along the ‘stair-case diagonal’ running from top-left to bottom-right of the reduced matrix.

```
1 U[1]
```

(0, 2)

And here is that column space (again, I will show each vector in the result separately):

```
1 C = A.columnspace()
2 C[0]
```

$$\begin{bmatrix} 1 \\ 2 \\ -1 \end{bmatrix}$$

```
1 C[1]
```

$$\begin{bmatrix} 3 \\ 9 \\ 3 \end{bmatrix}$$

In this case, all solutions to the original system of equations exist on a 2-dimensional (flat) plane that cuts across 3-dimensional space. The two vectors above define that plane. Every point on the plane is some linear combination of the above two vectors.

The ‘nullspace’ of the matrix is the set of vectors defining the sub-space that would be mapped onto zero by the matrix. All linear combinations of vectors in this space will be transformed onto zero by the matrix.

```
1 N = A.nullspace()
2 N[0]
```

$$\begin{bmatrix} -3 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

```
1 N[1]
```

$$\begin{bmatrix} 1 \\ 0 \\ -1 \\ 1 \end{bmatrix}$$

Which can be demonstrated with some examples:

```
1 A * N[0]
```

$$\begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

```
1 A * (2 * N[0])
```

$$\begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

```
1 A * (5 * N[0] + 7 * N[1])
```

$$\begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

Whereas all other vectors do not have this property:

```
1 NZ = Matrix([-4,1,0,0])
2 NZ
```

$$\begin{bmatrix} -4 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

```
1 A * NZ
```

$$\begin{bmatrix} -1 \\ -2 \\ 1 \end{bmatrix}$$

5.6 matrices in Numpy

The above should help you understand some of the operations of matrices and Linear Algebra. In practice however, sympy is much less frequently used in Machine Learning than Numpy. In general, in ML we need to be doing fast, numerical computation rather than slower, symbolic manipulation.

In this section I review the use of 'Numpy' - a very common library for numeric calculations in the field of Machine Learning.

```

1 import numpy as np
2
3 x = np.array([[1,2,2],[4,5,6]])
4 y = np.array([[7,8],[9,3],[2,1]])
5
6 print(x, "\n")
7 print(y)

```

```

[[1 2 2]
 [4 5 6]]

```

```

[[7 8]
 [9 3]
 [2 1]]

```

Note that numpy's data type we are using to represent Matrices is actually called an 'array'

Numpy provides a variety of functions for manipulating matrices, including:

5.6.1 Matrix Addition

Here I will define a new matrix 'p' to make the addition more interesting.

Note that addition is only define for matrices of the same 'shape'. I.e. with equivalent numbers of rows and columns.

Addition of matrices is particularly easy to read and write due to the re-use of the '+' operator:

```

1 p = np.array([[4,7,1],[4,3,2]])
2 print("x =\n", x)
3 print("\np =\n", p)
4 print("\nx + p =\n",x + p)

```

```

x =
[[1 2 2]
 [4 5 6]]

```

```

p =
[[4 7 1]
 [4 3 2]]

```

```

x + p =
[[5 9 3]
 [8 8 8]]

```

5.6.2 Matrix Subtraction

Analogously, the '-' operator performs element-by-element matrix (array) subtraction:

```
1 print("x =\n", x)
2 print("\np =\n", p)
3 print("\nx - p =\n", x - p)
```

```
x =
[[1 2 2]
 [4 5 6]]
```

```
p =
[[4 7 1]
 [4 3 2]]
```

```
x - p =
[[-3 -5 1]
 [ 0  2 4]]
```

5.6.3 Matrix multiplication by a Constant

If you simply want to multiply each element of a matrix by a constant, then that can be achieved using the "*" operator:

```
1 print (5 * x)
```

```
[[ 5 10 10]
 [20 25 30]]
```

Otherwise, if you want to multiply two matrices together, this can be achieved with the following. Notice, as the convention with matrix multiplication that the rank of the matrices has to match: the number of columns in x has to match the number of rows in y. Notice also the order of multiplication and the shape of the result:

```
1 print("x =\n", x)
2 print("\ny =\n", y)
3 print("\nmatmul(x,y) =")
4 print(np.matmul(x,y))
```

```

x =
[[1 2 2]
 [4 5 6]]

y =
[[7 8]
 [9 3]
 [2 1]]

matmul(x,y) =
[[29 16]
 [85 53]]

```

or .. using the '@' operator rather than the more familiar '*' operator for matrix multiplication

```

1 print(x @ y)

```

```

[[29 16]
 [85 53]]

```

5.6.4 Dot-product (Sometimes called 'inner product')

Numpy also provides a function to calculate the 'dot-product' between two 1-dimensional arrays (also frequently referred to as 'vectors'). The result of the numpy dot-product between two vectors is a number (scaler).

The dot-product measures the 'projection' of one vector onto another. That is, it provides a measure of similarity in the 'direction' of vectors. If vectors are orthogonal (at right-angles) to each other, then the scaler product is defined as zero. Otherwise, the scaler product indicates the size of the 'shadow' one vector would cast on the other.

In text books this operation is often shown as $x.y^T$

When thinking about vectors in space it is more frequently indicated as $|x|.|y|.cos(theta)$, where theta is the angle between the lines.

Both of these provide the same result.

The scaler product can be calculated in numpy using 'dot':

```

1 v1 = np.array([1,2,3])
2 v2 = np.array([4,5,6])
3 print (np.dot(v1,v2))

```


Finally, numpy also provides an 'inner' function. For 1-D vectors the result of 'inner' is identical to the 'dot' product. In higher dimensions these operators perform different operations - however those are beyond the scope of this course.

```
1 print (np.inner(v1,v2))
```

32

5.6.5 Matrix Division by a constant

We divide each element of a matrix by a constant like this:

```
1 print(np.divide(x,2))
```

```
[[0.5  1.   1.  ]
 [2.   2.5  3.  ]]
```

By analogy with the above, you can also use the '/' operator to divide a matrix by a constant:

```
1 print(x/2)
```

```
[[0.5  1.   1.  ]
 [2.   2.5  3.  ]]
```

We can do an element-by-element divide between two matrices of the same shape:

```
1 print (np.divide(x,p))
```

```
[[0.25      0.28571429  2.          ]
 [1.        1.66666667  3.          ]]
```

And again, use the '/' operator if more convenient:

```
1 print (x/p)
```

```
[[0.25      0.28571429  2.          ]
 [1.        1.66666667  3.          ]]
```

5.6.6 Determinant of a matrix

The Determinant of a matrix is a measure of the ‘scaling’ that the matrix generates when it is applied to a series of vectors. This is very well explained here:

<https://towardsdatascience.com/what-really-is-a-matrix-determinant-89c09884164c>

Determinants can be calculated for square matrices.

Numpy provides an easy-to-use function to calculate determinants:

```
1 s = np.array([[1,2,3], [14,25,36], [74,38,29]])
2 print(np.linalg.det(s))
```

```
-81.000000000000024
```

5.6.7 Inverse of a matrix

Finally, numpy provides a function to determine the inverse of a matrix:

```
1 s_inv = np.linalg.inv(s)
2 print(s_inv)
```

```
[[ 7.9382716 -0.69135802  0.03703704]
 [-27.87654321  2.38271605 -0.07407407]
 [ 16.27160494 -1.35802469  0.03703704]]
```

matrices, of course, have the property that if they are multiplied by their inverse the result is the Identity matrix (1's on the diagonal, 0's elsewhere):

```
1 print( np.matmul(s,s_inv))
```

```
[[ 1.00000000e+00  6.66133815e-16  6.93889390e-18]
 [-2.84217094e-14  1.00000000e+00  1.38777878e-16]
 [-7.10542736e-15  7.32747196e-15  1.00000000e+00]]
```

Which is a close approximation to:

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

5.6.8 Solutions to Linear Equations

As mentioned in the first section, series of linear equations can be represented as matrices.

$$1x + 2y = 1$$

$$3x + 5y = 2$$

In the context of Machine Learning you will often see such equations represented in the form of matrices:

$$\begin{bmatrix} 1 & 2 \\ 3 & 5 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$$

Such systems of linear equations may be soluble using numpy. By ‘soluble’ I mean that we can find values for x and y for which the system of equations is consistent.

```
1 M1 = np.array([[1, 2], [3, 5]])
2
3 M2 = np.array([1, 2])
4
5 A = np.linalg.solve(M1, M2)
6
7 print(A)
```

```
[-1.  1.]
```

In this case, substituting a value of ‘-1’ for x, and ‘1’ for y, produces a constant set of equations:

```
1 print(1 * -1 + 2 * 1)
```

```
1
```

```
1 print (3 * -1 + 5 * 1)
```

```
2
```

Of course, in some cases there will be no solution to a set of linear equation. For example:

$$1x + 2y = 1$$

$$2x + 4y = 3$$

represented as:

$$\begin{bmatrix} 1 & 2 \\ 2 & 4 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} 1 \\ 3 \end{bmatrix}$$

The following should produce a ‘Singular Matrix’ error, since there are no values for which the equations hold true:

```
1 M1 = np.array([[1, 2], [2, 4]])
2
3 M2 = np.array([1, 3])
4
5 A = np.linalg.solve(M1, M2)
6
7 print(A)
```

LinAlgError: Singular matrix

5.6.9 Matrix (or vector) Norm

It is often useful to be able to compute the ‘length’ of a vector. There are many ways to do this (technically there are an infinite number of ways of doing it!). However, the most common methods for computing lengths (distances) in Machine learning are the L1 and L2 norms.

The L2 norm is what, in standard geometry we call a “Euclidean distance” (or “Euclidean norm”) and is simply the square root or the sum of the squares of the dimensions. Here is the simplest example .. a 3,4,5 triangle:

```
1 myVector = [3,4]
2 print(np.linalg.norm(myVector))
```

5.0

Numpy extends this into higher dimensions, although calculated in an exactly analogous manner:

```
1 myVector = [1,2,3]
2 print(np.linalg.norm(myVector))
```

3.7416573867739413

Numpy can also calculate the ‘L1’ norm as follows. This is sometimes referred to as the ‘Manhattan distance’. The L1 is simply the sum of each vector dimension:

```
1 print(np.linalg.norm(myVector, ord = 1))
```

6.0